

Homework #7

Winter 2020, STATS 509

Dino Bektesevic

Problem 1: Sampling Distributions

A researcher plans to carry out an opinion survey regarding a yes/no question. Suppose that (unknown to the researcher) 89% of people in the target population answer ‘yes’. We will use 1 to denote ‘yes’ and 0 to denote ‘no’.

- a. Let X_1 be the response from the first person the researcher asks. Before this person is asked, what is the distribution of X_1 . What is $E[X_1]$ and $V(X_1)$?

The researcher plans to obtain a random sample of size 25: X_1, \dots, X_{25} .

(Suppose either that the researcher is sampling with replacement or that the population is so large that we may regard the samples as being taken with replacement.)

- b. Write down the distribution of X_i for $i = 1, \dots, 25$; also write down $E[X_i]$, $V(X_i)$ and $\text{Cov}(X_i, X_j)$ for $i \neq j$.

- c. What is the distribution of \bar{X} ? *Hint: see lecture notes.*

What is the mean and variance of this distribution?

- d. Use R or Python, replicate the experiment 10,000 times.

Hint: Recall that a Bernoulli(p) random variable is just a Binomial(1, p) random variable; thus to obtain a sample of size 25 from a Bernoulli(0.89) random variable, one might use:

```
p <- 0.89
```

```
n <- 25
```

```
x <- rbinom(n,1,p)
```

The vector \mathbf{x} will contain the 25 Bernoulli(p) observations, from which you can compute the sample mean. This should then be put in a loop which repeats this process, 10K times, storing the result. Construct a histogram with the distribution of \bar{X} for each sample. Does the distribution of sample means appear to be normal? Explain your answer.

- e. Find the mean and variance of \bar{X} , based on your 10K simulations. Does this agree with your calculation in c.?

- f. Compute the mean squared error of \bar{X} as an estimate of the unknown proportion $\theta = 0.89$ (i.e. find the squared error of the estimate, which is $(\bar{X} - \theta)^2$, in each repetition and then average over all repetitions). What do you notice? Give a simple explanation by referring to a result from the course.

Problem 2

Given a sample X_1, \dots, X_n of independent $\text{Poisson}(\lambda)$ random variables, a researcher intends to use \bar{X} , the sample mean, as an estimate of λ . Suppose that $n = 20$, and $\lambda = 4$.

- a. Write down the mean and variance of this distribution.
- b. Using R or Python, replicate the experiment 10,000 times. Construct a histogram with the distribution of \bar{X} for each sample. Does the distribution of sample means appear to be normal? Explain your answer.
- c. Find the mean and variance of \bar{X} , based on your 10K simulations. Does this agree with your calculation in a.?
- d. Compute the mean squared error of \bar{X} as an estimate of λ . Again what do you notice?

Problem 3: Hypothesis Tests

Suppose that Y_1, \dots, Y_n are i.i.d samples from a Poisson distribution with parameter λ .

- a. Find the log of the likelihood: $\ln f(y_1, \dots, y_n | \lambda)$; Poisson distribution is given as

$$f(y_i | \lambda) = \frac{e^{-\lambda} \lambda^{y_i}}{y_i!}$$

The likelihood can be written as:

$$L(\lambda | y_1 \dots y_n) = \prod_{i=1}^n f(y_i | \lambda) = \frac{\lambda^{\sum_{i=1}^n y_i} e^{-n\lambda}}{\prod_{i=1}^n y_i!}$$

Log likelihood is then:

$$l(\lambda | y_1 \dots y_n) = \ln(\lambda) \sum_{i=1}^n y_i - n\lambda - \ln \prod_{i=1}^n y_i!$$

- b. By differentiating the log likelihood, find the value $\hat{\lambda}$ of λ that maximizes the likelihood; confirm that $\hat{\lambda}$ is a maximum. The estimate $\hat{\lambda}$ is called the *maximum likelihood estimator (MLE)*.

$$\begin{aligned} \frac{d}{d\lambda} l(\lambda | y_1 \dots y_n) &= \frac{d}{d\lambda} \ln(\lambda) \sum_{i=1}^n y_i - \frac{d}{d\lambda} n\lambda - \frac{d}{d\lambda} \ln \prod_{i=1}^n y_i! = 0 \\ 0 &= \frac{\sum_{i=1}^n y_i}{\hat{\lambda}} - n \\ n\hat{\lambda} &= \sum_{i=1}^n y_i \\ \hat{\lambda} &= \frac{\sum_{i=1}^n y_i}{n} \\ \hat{\lambda} &= \bar{y} \end{aligned}$$

To confirm it's a maximum we have to take a look at the second derivative:

$$\begin{aligned} \frac{d^2}{d\lambda^2} l(\lambda | y_1 \dots y_n) &= \frac{d}{d\lambda} \frac{\sum_{i=1}^n y_i}{\lambda} \\ &= -\frac{\sum_{i=1}^n y_i}{\lambda^2} \end{aligned}$$

Since it's always negative this is a maximum.

Suppose that we wish to test $H_0: \lambda = \lambda_0$ vs. $H_1: \lambda \neq \lambda_0$.

- c. Using your answer to b. write down the generalized likelihood ratio test statistic (LRT).

Hint: Your answer should be a function of n , λ_0 and $\hat{\lambda}$.

$$\begin{aligned} \lambda_{lrt}(y_1 \dots y_n) &= \frac{f(y_1 \dots y_n | \lambda_0)}{\sup_{\lambda \in \mathbb{R}} f(y_1 \dots y_n | \lambda)} \\ &= \frac{\lambda_0^{\sum_{i=1}^n y_i} e^{-n\lambda_0}}{\hat{\lambda}^{\sum_{i=1}^n y_i} e^{-n\hat{\lambda}}} \\ &= \frac{\lambda_0^{\sum_{i=1}^n y_i} e^{-n\lambda_0}}{\bar{y}^{\sum_{i=1}^n y_i} e^{-n\bar{y}}} \\ &= e^{-n(\lambda_0 - \bar{y})} \left(\frac{\lambda_0}{\bar{y}} \right)^{n\bar{y}} \end{aligned}$$

d. Consider Bortkiewicz's Prussian Cavalry horse-kick fatality data:

No. of fatalities	0	1	2	3	4
No. of years	109	65	22	3	1

The original table from Bortkiewicz's 1898 book *Das Gesetz der kleinen Zahlen* is: <https://archive.org/stream/dasgesetzderklei00bortrich#page/n66/mode/1up>

Find the value of $\sum_{i=1}^n y_i$, and use your answer to b. to find the value of the MLE $\hat{\lambda}$.

Hint: The table contains data on $n = 200$ observations, y_1, \dots, y_{200} , of which 109 were 0; there were 65 which were 1 and so on.

$$\sum_{i=1}^{200} y_i = 0 \cdot 109 + 65 + 2 \cdot 22 + 9 + 4 = 122$$

So the $\bar{y} = \frac{122}{200} \approx 0.6$

e. Use your answer to d. to find the LRT statistic for the hypothesis test with $\lambda_0 = 1$.

$$\lambda_{lrt}(y_1 \dots y_n) = e^{-n(\lambda_0 - \bar{y})} \left(\frac{\lambda_0}{\bar{y}} \right)^{n\bar{y}} \quad (1)$$

$$= e^{-200(1-0.6)} \left(\frac{200}{122} \right)^{200 \frac{122}{200}} \quad (2)$$

$$= e^{-80} \left(\frac{200}{122} \right)^{122} \quad (3)$$

$$= 2.79 \cdot 10^{-9} \quad (4)$$

f Report the approximate p-value.

Hint: calculate $-2\log(LRT)$, and compare to the appropriate χ^2 distribution. See Lecture 10, slide 43. Recall that small values of the LRT correspond to evidence against H_0 .

Problem 4

Suppose that we are planning an experiment to test hypotheses about the mean of a population that is known to be normal with standard deviation $\sigma = 4$. We wish to test the null hypothesis $H_0 : \mu = 0$ vs. the alternative $H_A : \mu > 0$. We intend to use a likelihood ratio test with significance level $\alpha = 0.05$.

- a. For which values of \bar{X} , the sample mean, will we reject the null hypothesis. Express your answer as a function of sample size, n .

We are given that $\bar{X} \sim N(\mu, \sigma)$. Rewriting this to standard normal yields:

$$Y = \frac{\sqrt{n}(\bar{X} - \mu)}{\sigma} \sim N(0, 1)$$

Under the null hypothesis $\mu = 0$ so we relate α to c by:

$$\begin{aligned} \alpha &= P(\bar{X} > c | \mu = 0) \\ &= P\left(\frac{\sqrt{n}(\bar{X} - \mu)}{\sigma} > \frac{\sqrt{n}(c - \mu)}{\sigma} \middle| \mu = 0\right) \\ &= P\left(Y > \frac{\sqrt{n}(c - \mu)}{\sigma} \middle| \mu = 0\right) \\ &= 1 - P\left(Y < \frac{\sqrt{n}(c - \mu)}{\sigma} \middle| \mu = 0\right) \\ 1 - \alpha &= \Phi\left(\frac{\sqrt{n}(c - \mu)}{\sigma} \middle| \mu = 0\right) \\ 0.95 &= \Phi\left(\sqrt{n}\frac{c - \mu}{\sigma} \middle| \mu = 0\right) \end{aligned}$$

The CDF of a standard normal distribution equals 0.95 at the standard score of 1.644854 so we can determine c to be:

$$\begin{aligned} \sqrt{n}\frac{c}{4} &= 1.644854 \\ c &= \frac{6.579416}{\sqrt{n}} \end{aligned}$$

- b. Suppose that we plan to obtain a sample of size 36. The researcher thinks that if the alternative is true then perhaps $\mu = 0.3$. Calculate the power of the test to reject the null hypothesis under this particular alternative hypothesis.

Equivalently to the previous problem we can start by writing:

$$\begin{aligned} 1 - \beta &= 1 - P\left(Y < \frac{\sqrt{n}(c - \mu)}{\sigma} \middle| \mu = 0.3\right) \\ &= 1 - \Phi\left(\frac{\sqrt{n}(c - \mu)}{\sigma}\right) \\ &= 1 - \Phi\left(\frac{\sqrt{n}(c - \mu)}{\sigma} \middle| \mu = 0.3\right) \\ &= 1 - \Phi\left(\frac{\sqrt{n}\left(\frac{6.579416}{\sqrt{n}} - \mu\right)}{\sigma} \middle| \mu = 0.3\right) \\ &= 1 - \Phi\left(\frac{6\left(\frac{6.579416}{6} - 0.3\right)}{4}\right) \\ &= 1 - \Phi(1.194854) \\ &= 0.116 \end{aligned}$$

- c. Continuing from b., the scientist who is planning the experiment wishes to have power at least 90%. (Thus your calculation in b. shows that more than 36 observations are required.) Find approximately the smallest sample size at which this power can be achieved, against the specific alternative $\mu = 0.3$.

Hint: repeat the calculation performed in b. at different sample sizes using trial and error: you may wish to use R or a spreadsheet to speed up this calculation.

- d. Suppose that the researcher obtains 200 samples, and observes $\bar{x} = 0.65$. Compute the p-value for this hypothesis test.

$$\begin{aligned} p_v &= P(\bar{X} > \bar{x} | H_0) \\ &= 1 - \Phi\left(\frac{\sqrt{200}(0.65 - 0)}{4}\right) \\ &= 0.0108 \end{aligned}$$

Problem 5

A researcher performs a sequence of independent experiments, up to and including the first ‘success’, after which the researcher stops. Each experiment has the same probability p of success. Let T be the number of experiments performed (including the first observed success). The researcher wishes to test the null hypothesis

- a. $H_0: p = 0.25$,
- b. against the alternative hypothesis
- c. $H_1: p > 0.25$.

The researcher proposes to reject the null hypothesis if $T < 4$.

- a. What is the significance level α of the test proposed by the researcher?

$$\begin{aligned}
 \alpha &= P(T < 4 | H_0) \\
 &= P(T \leq 3 | H_0) \\
 &= \sum_{t=1}^3 p_0(1-p_0)^{t-1} \\
 &= \left(\frac{1}{4} + \frac{3}{16} + \frac{9}{64} \right) \\
 &= \frac{37}{64}
 \end{aligned}$$

- b. What is the power of the researcher’s test against the specific alternative hypothesis that $p = 0.5$.

$$\begin{aligned}
 1 - \beta &= P(T < 4 | H_1) \\
 &= P(T \leq 3 | H_1) \\
 &= \sum_{t=1}^3 p_1(1-p_1)^{t-1} \\
 &= \sum_{t=1}^3 0.5^t \\
 &= \left(\frac{1}{2} + \frac{1}{4} + \frac{1}{8} \right) \\
 &= \frac{7}{8}
 \end{aligned}$$

- c. Re-express the researcher’s rule for rejecting the null hypothesis in terms of the likelihood ratio:

$$\text{LRT} = p(t \mid p = 0.25) / p(t \mid p = 0.5).$$

Specifically, find the value ℓ such that the researcher will reject $H_0: p = 0.25$ in favor of $p = 0.5$ if $\text{LRT} < \ell$.

(Note: There will be a range of values for ℓ that will give the same test.)

$$\begin{aligned}
 \Lambda(t) &= \frac{P(T = t | H_0)}{P(T = t | H_1)} \\
 &= \frac{p_0(1-p_0)^{t-1}}{p_1(1-p_1)^{t-1}} \\
 &= \frac{1}{2} \left(\frac{3}{2} \right)^{t-1}
 \end{aligned}$$

Hint: (For all parts) Geometric distribution!

Problem 6: Confidence Intervals

Consider a 95% confidence interval for the mean height μ in a population. Which of the following are true or false:

- a. Before taking our sample, the probability of the resulting 95% confidence interval containing μ is 0.95.
True.
- b. If we take a sample and compute a 95% confidence interval for μ to be $[1.2, 3.7]$ then $P(\mu \in [1.2, 3.7]) = 0.95$.
False. We can only think of μ as being random before taking a sample. Afterwards, it's just a number.
- c. Before taking our sample, the center of a 95% confidence interval for the population mean is a random variable.
True, before taking a sample the center of our interval is a random variable because it's based on the sampling a random variable.
- d. 95% of individuals in the population have heights that lie in the 95% confidence interval for μ .
False. CI tells us the confidence that after drawing a random sample it's mean will lie in that range, not the other way around.
- e. Over hypothetical replications out of one hundred 95% confidence intervals for μ , on average 95 will contain μ .
True.
- f. After obtaining our sample, the resulting confidence interval either does or does not contain μ .
True, we don't know if our sampled CI actually contains or doesn't contain the mean.

Problem 7

Goldberger Qu. 11.6 (Assume that the observations are drawn from a Normal Distribution and see Lecture 10, slide 46.)